

Operating System Principles

操作系统原理

Process/Thread Scheduling

李旭东

leexudong@nankai.edu.cn

Nankai University




Objectives

- Scheduler
- Process Behavior
- Scheduling Mode
- Scheduling Criteria
- Scheduling Algorithms
- Thread Scheduling



MultiProgramming

- Scheduler
- Scheduling algorithm

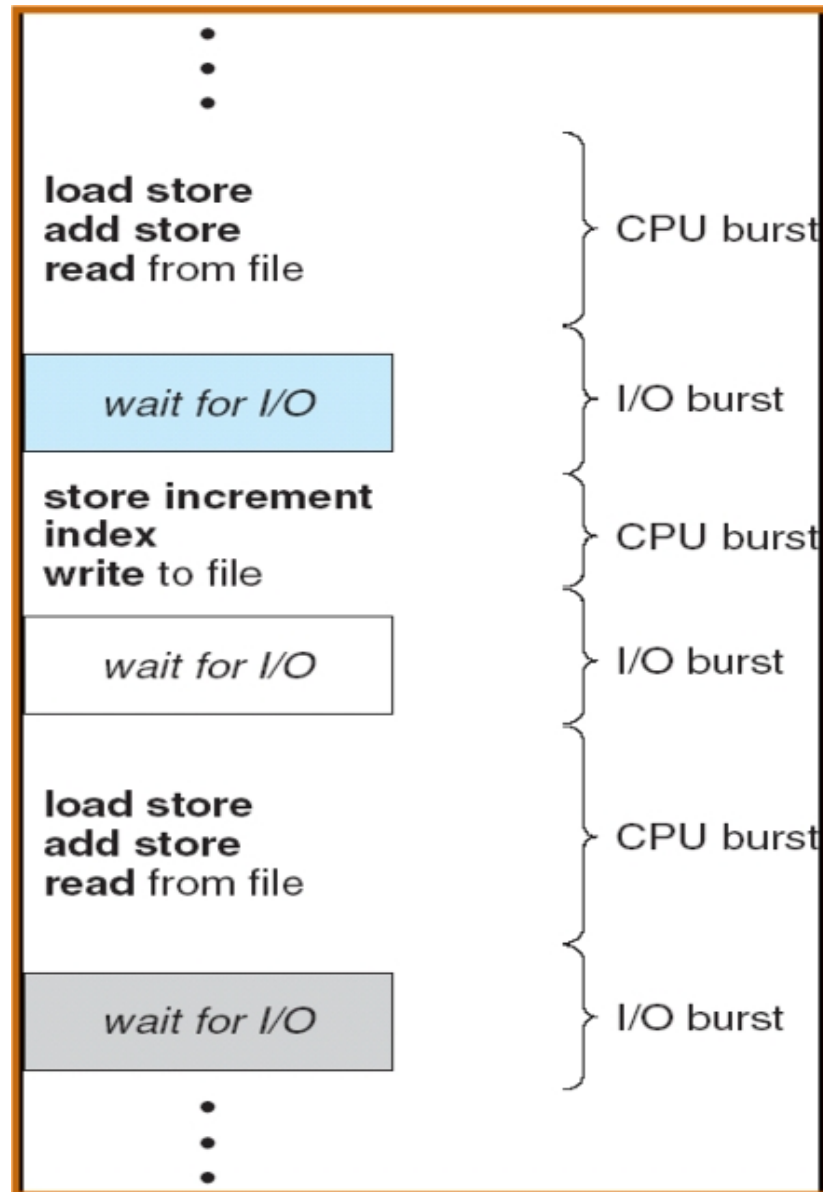




Scheduler


- Short-term Scheduler
 - CPU
- Middle-term Scheduler
 - Memory
- Long-term Scheduler
 - Job

CPU And I/O Bursts in a Process




Process Behavior

- Compute-bound
- I/O-bound



Multi-Processes Trace





When to Schedule

- A new process is created
- A process exits
- A process blocks on I/O, on a semaphore, or for some other reason
- An I/O interrupt occurs



Dispatcher

- A module that gives control of the CPU to the process selected by the short-term scheduler
 - Switching context
 - Switching to user mode
 - Jumping to the proper location in the user program to restart that program
- Dispatch latency 调度延迟
 - The time it takes for the dispatcher to stop one process and start another running



Scheduling Modes

- Preemptive
 - 抢占式
- Nonpreemptive
 - 非抢占式，非剥夺式



Categories of Scheduling Algorithms

- Batch
- Interactive
- Real-time



Scheduling Criteria

- CPU utilization 利用率
- Throughout 吞吐量
- Turnaround time 周转时间
 - Waiting to get into memory
 - Waiting in the ready queue
 - Executing on the CPU
 - Doing I/O
- Waiting time
- Response time
- ...



Scheduling Algorithm Goals

All systems

Fairness - giving each process a fair share of the CPU

Policy enforcement - seeing that stated policy is carried out

Balance - keeping all parts of the system busy

Batch systems

Throughput - maximize jobs per hour

Turnaround time - minimize time between submission and termination

CPU utilization - keep the CPU busy all the time

Interactive systems

Response time - respond to requests quickly

Proportionality - meet users' expectations

Real-time systems

Meeting deadlines - avoid losing data

Predictability - avoid quality degradation in multimedia systems



Scheduling in Batch System

- First-come first-served
- Shortest job first
- Shortest remaining Time next

First-come first-served

- Average waiting time

<u>Process</u>	<u>Burst Time</u>
----------------	-------------------

P_1	24
-------	----

P_2	3
-------	---

P_3	3
-------	---




$$Awt = (0 + 24 + 27) / 3 = 17$$




$$Awt = ?$$

Shortest job first



(a)



(b)

1. nonpreemptive
2. preemptive

Figure 2-40. An example of shortest job first scheduling.

(a) Running four jobs in the original order.

(b) Running them in shortest job first order.



Quiz

Process	Burst Time
---------	------------

P_1	6
-------	---

P_2	8
-------	---

P_3	7
-------	---

P_4	3
-------	---

SJF: AWT=?

FCFS: AWT=?



Shortest job first

- How to predict length of the next CPU burst?
 - exponential average

$$T_{n+1} = at_n + (1-a)T_{n-1}, \quad 0 \leq a \leq 1$$

t_n the length of the n th CPU burst

T_{n+1} the predicted value of the next CPU burst

$$T_{n+1} = at_n + (1-a)at_{n-1} + \cdots + (1-a)^j at_{n-j} + \cdots + (1-a)^{n+1}T_0$$



Shortest remaining Time next

- i.e. Preemptive SJF scheduling

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

Nonpreemptive SJF scheduling: AWT=?

Preemptive SJF scheduling: AWT=?



Scheduling in Interactive System

- Round-Robin Scheduling
- Priority Scheduling
- Multiple Queues
- Shortest Process Next
- Guaranteed Scheduling
- Lottery Scheduling
- Fair-Share Scheduling

Round-Robin Scheduling





Figure 2-41. Round-robin scheduling.

(a) The list of runnable processes.


(b) The list of runnable processes after B uses up its quantum.



Round-Robin Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3


a time quantum of 4 milliseconds



$$\text{AWT} = (6 + 4 + 7) / 3 = 5.66$$

Quantum Value

- How a smaller time quantum increases context switches



Priority Scheduling

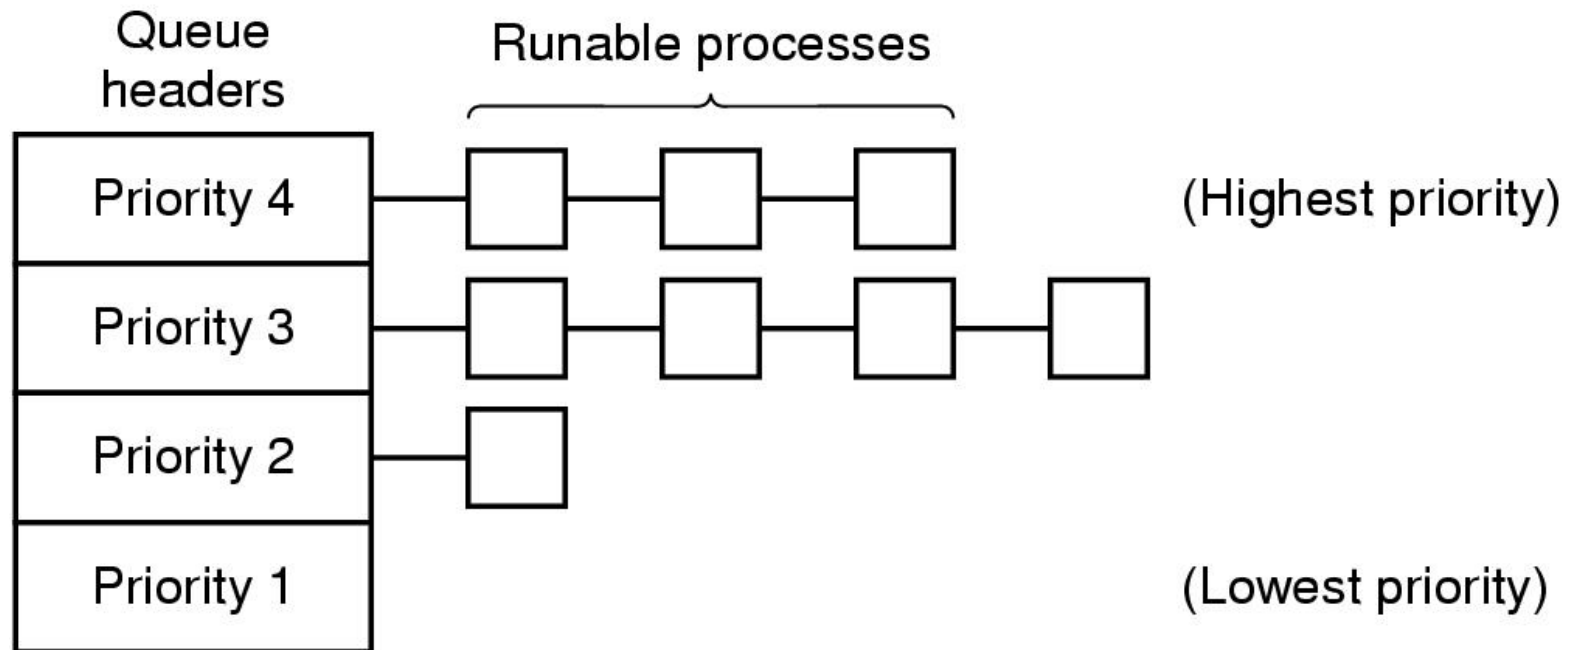



Figure. A scheduling algorithm with four priority classes.

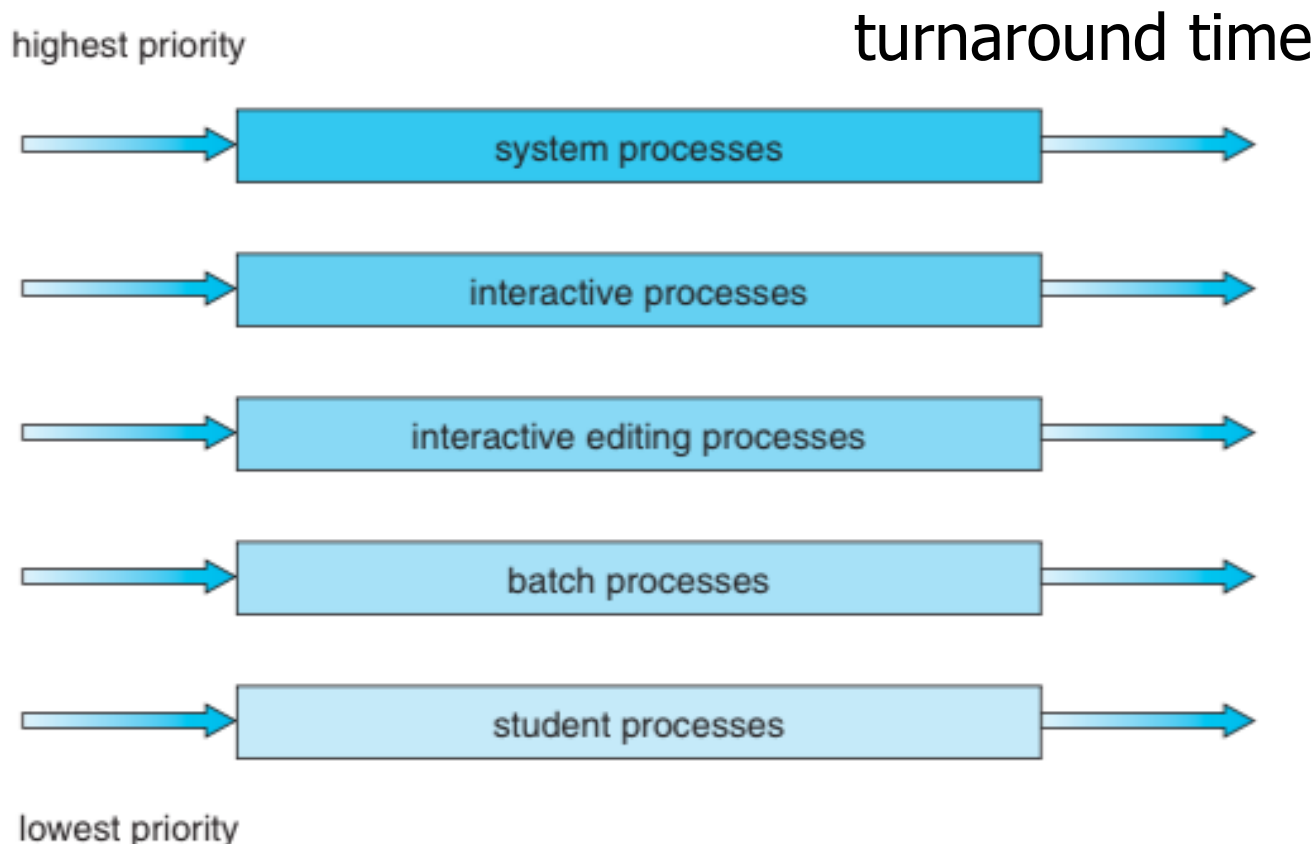
Multilevel Feedback Queue Scheduling

- Idea
 - Separate processes with different CPU-burst characteristics
 - Allow a process to move between queues



Multilevel Queue Scheduling

- **Foreground** (interactive) processes
- **Background** (batch) processes





Mutli-level Feedback Queue Scheduling

- the scheduler is defined by the following parameters:
 - The number of queues
 - The scheduling algorithm for each queue
 - The method used to determine when to upgrade a process to a higher-priority queue
 - The method used to determine when to demote a process to a lower-priority queue
 - The method used to determine which queue a process will enter when that process needs service

Lottery Scheduling





More Scheduling Algorithms

- Guaranteed Scheduling
- Fair-Share Scheduling
- ...

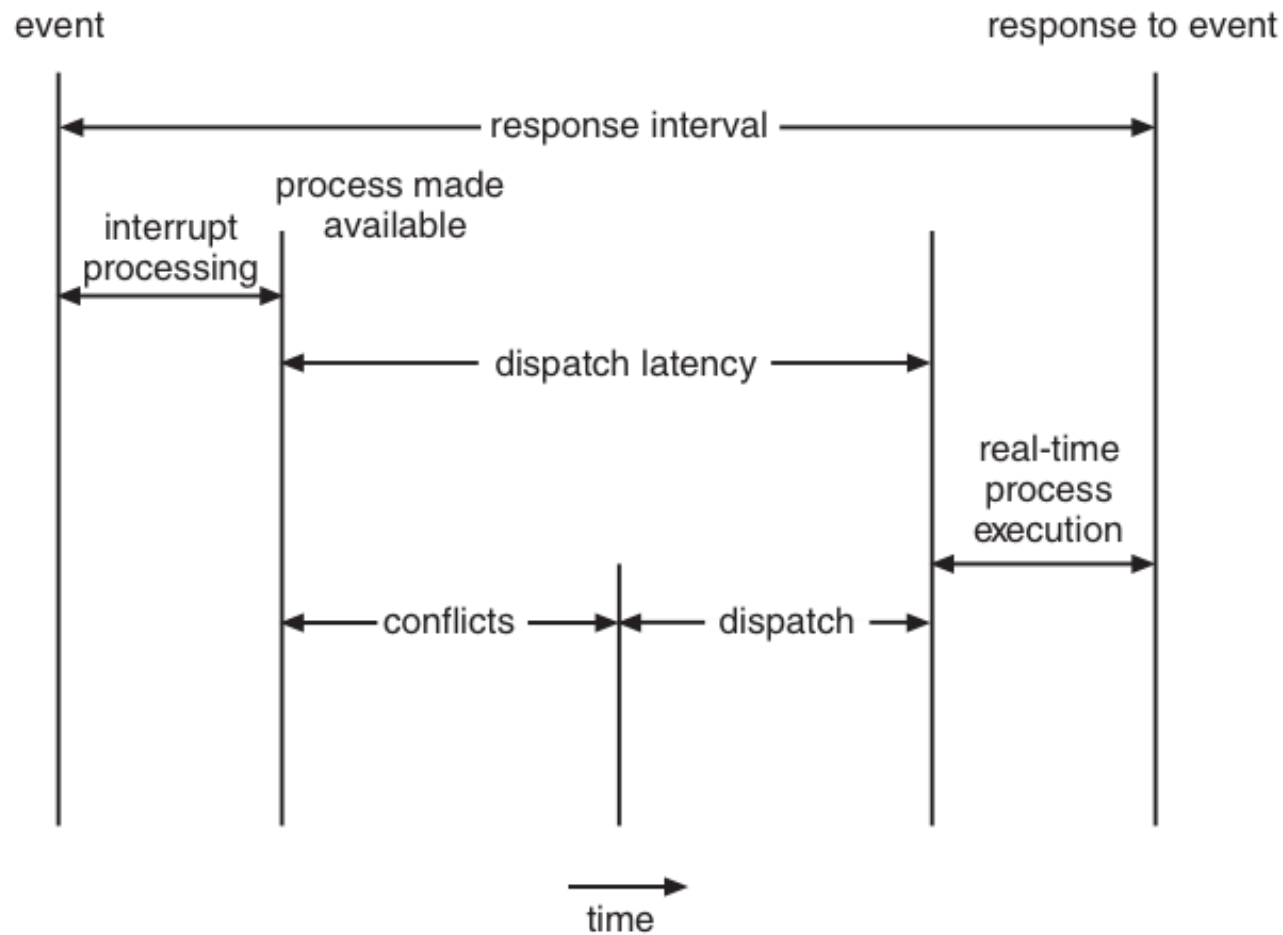


Scheduling in Real-time System

- Categories I
 - Hard real time
 - Soft real time
- Categories II
 - Periodic
 - Aperiodic
- Categories III
 - Static
 - dynamic


Real-time CPU Scheduling

- Minimizing Latency 最小化延迟
 - Interrupt latency, Dispatch latency




Real-time CPU Scheduling

- Priority-Based Scheduling



Example: P2 has a higher priority than P1




Real-time CPU Scheduling

- Rate-Monotonic Scheduling
 - 单一速率
 - a static priority policy with preemption

Example:


P1 a higher priority than P2;

the period of P1 is shorter than that of P2




Real-time CPU Scheduling

- Earliest-Deadline-First Scheduling
 - dynamically assigns priorities according to deadline



Missing deadlines with rate-monotonic scheduling



Earliest-deadline-first scheduling



Real-time CPU Scheduling

- Proportional 成比例的 Share Scheduling
 - Proportional share schedulers operate by allocating T shares among all applications
 - An application can receive N shares of time, thus ensuring that the application will have N/T of the total processor time



Real-time CPU Scheduling

- POSIX Real-Time Scheduling
 - SCHED_FIFO
 - SCHED_RR
 - SCHED_OTHER
 - `pthread_attr_getsched_policy(pthread_attr_t *attr, int *policy)`
 - `pthread_attr_setsched_policy(pthread_attr_t *attr, int policy)`



Multiple-Processor Scheduling

- multiple CPUs
 - load sharing becomes possible—but scheduling problems become correspondingly more complex
- Approaches to Multiple-Processor Scheduling
 - asymmetric multiprocessing
 - all scheduling decisions, I/O processing, and other system activities handled by a single processor—the master server.
 - The other processors execute only user code.
 - symmetric multiprocessing (SMP)
 - Each processor is self-scheduling. All processes may be in a common ready queue, or each processor may have its own private queue of ready processes.



Multiple-Processor Scheduling

- Processor Affinity 处理器亲和性
 - Consider what happens to cache memory when a process has been running on a specific processor. The data most recently accessed by the process populate the cache for the processor.
 - As a result, successive memory accesses by the process are often satisfied in cache memory.
- Deference Forms of Processor Affinity
 - soft affinity
 - When an operating system has a policy of attempting to keep a process running on the same processor—but not guaranteeing that it will do so
 - hard affinity
 - `sched_setaffinity()` system call



Linux Scheduler

- /kernel/sched/core.c

- static void __sched notrace __schedule(bool preempt)
- static __always_inline struct rq *
context_switch(struct rq *rq, struct task_struct *prev,
struct task_struct *next, struct rq_flags *rf)
- static inline struct task_struct *
pick_next_task(struct rq *rq, struct task_struct *prev,
struct rq_flags *rf)

Linux Scheduler (cont.,)

Filter tags

v4

v4.14

v4.14-rc8

v4.14-rc7

v4.14-rc6

v4.14-rc5

v4.14-rc4

v4.14-rc3

v4.14-rc2

v4.14-rc1

v4.13

v4.12

v4.11

v4.10

v4.9

v4.8

v4.7

v4.6

v4.5

v4.4

v4.3

v4.2

v4.1

v4.0

v3

v2.6

/ kernel / sched / core.c

Search Identifier

3193 }
3194
3195 /*
3196 * Pick up the highest-prio task:
3197 */
3198 static inline struct task_struct *
3199 pick_next_task(struct rq *rq, struct task_struct *prev, struct rq_flags *rf)
3200 {
3201 const struct sched_class *class;
3202 struct task_struct *p;
3203
3204 /*
3205 * Optimization: we know that if all tasks are in the fair class we can
3206 * call that function directly, but only if the @prev task wasn't of a
3207 * higher scheduling class, because otherwise those loose the
3208 * opportunity to pull in more work from other CPUs.
3209 */
3210 if (likely((prev->sched_class == &idle_sched_class ||
3211 prev->sched_class == &fair_sched_class) &&
3212 rq->nr_running == rq->cfs.h_nr_running)) {
3213
3214 p = fair_sched_class.pick_next_task(rq, prev, rf);
3215 if (unlikely(p == RETRY_TASK))
3216 goto again;
3217
3218 /* Assumes fair_sched_class->next == idle_sched_class */
3219 if (unlikely(!p))
3220 p = idle_sched_class.pick_next_task(rq, prev, rf);
3221
3222 return p;
3223 }
3224
3225 again:
3226 for_each_class(class) {
3227 p = class->pick_next_task(rq, prev, rf);
3228 if (p) {
3229 if (unlikely(p == RETRY_TASK))
3230 goto again;
3231 return p;
3232 }
3233 }
3234
3235 /* The idle class should always have a runnable task: */
3236 BUG();
3237 }
3238
3239 /*
3240 * __schedule() is the main scheduler function.

linux v4.14

powered by Elixir 0.2




Thread Scheduling

- Two levels of parallelism
- Thread scheduler
 - User-level thread
 - Kernel-level thread
 - (Hyper Thread)

Thread Scheduling

User-level thread: **process-contention scope** (PCS)
进程范围竞争




Possible: A1, A2, A3, A1, A2, A3

Not possible: A1, B1, A2, B2, A3, B3

(a) Possible scheduling of user-level threads with a 50-msec process quantum and threads that run 5 msec per CPU burst.

Thread Scheduling

Kernel-level thread: **system-contention scope** (SCS)



Possible: A1, A2, A3, A1, A2, A3

Also possible: A1, B1, A2, B2, A3, B3

(b) Possible scheduling of kernel-level threads with the same characteristics as (a).



Thread Scheduling Case

```
#include <pthread.h> #include <stdio.h>

#define NUM_THREADS 5

int main(int argc, char *argv[]){
    int i;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_t init(&attr);
    /* set the scheduling algorithm to PROCESS or SYSTEM */
    pthread_attr_t setscope(&attr, PTHREAD_SCOPE_SYSTEM);
    /* set the scheduling policy - FIFO, RT, or OTHER */
    pthread_attr_t setschedpolicy(&attr, SCHED_OTHER);
    /* create the threads */
    for (i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], &attr, runner, NULL);
```



Thread Scheduling Case

```
/* now join on each thread */
```

```
for (i = 0; i < NUM THREADS; i++)
```

```
    pthread_join(tid[i], NULL);
```

```
}
```

```
/* Each thread will begin control in this function */
```

```
void *runner(void *param)
```

```
{
```

```
    printf("I am a thread\n");
```

```
    pthread_exit(0);
```

```
}
```



Thread Scheduling

- Pthread Scheduling
 - PTHREAD_SCOPE_PROCESS schedules threads using PCS scheduling
 - PTHREAD_SCOPE_SYSTEM schedules threads using SCS scheduling
- pthread_attr_t setscope(pthread_attr_t *attr, int scope)
- pthread_attr_t getscope(pthread_attr_t *attr, int *scope)

Policy v.s. Mechanism

- Scheduling mechanism 调度机制
- Scheduling policy 调度策略






Summary

- Scheduler
- Process Behavior
- Scheduling Mode
- Scheduling Criteria
- Scheduling Algorithms
- Thread Scheduling
- ...



Q&A?



A word cloud of the word "thank you" in various languages and scripts. The words are arranged in a circular pattern, with "thank you" being the largest and most central. Other prominent words include "danke", "gracias", "merci", "obrigado", "dank je", "teşekkür ederim", "sukriya", "kop khun krap", "arigatō", "tak", "dakujem", "merci", "go raibh maith agat", "mochchakkeram", "tapadh leat", "hvala", "dziękuję", "bedankt", "nanni", "nandri", "kiitos", "dankie", "faafetai lava", "mauriuru", "koszönöm", "bayarlalaa", "gracie", "hvala", "sagolun", "didi madloba", "kam sah hamnida", "rahmat", "najes tuke", "chhorakaloutioun", "gratias ago", "gracies", "sulpáy", "go raibh maith agat", "mochchakkeram", "mamnun", "trugarez", "merci", "shukriya", "dhanyavadagal", "diolch", "eucharistw", "xiexie", "감사합니다", "তোমাকে ধন্যবাদ", "merci", "dakujem", "arigatō", "tak", "dakujem", "merci", "go raibh maith agat", "mochchakkeram", "mamnun", "trugarez", "merci", "shukriya", "dhanyavadagal", "diolch", "eucharistw", "xiexie", "감사합니다", "তোমাকে ধন্যবাদ", "merci", "dakujem", "arigatō", "tak", "dakujem", "merci", "go raibh maith agat", "mochchakkeram", "mamnun", "trugarez", "merci", "shukriya", "dhanyavadagal", "diolch", "eucharistw", "xiexie", "감사합니다", "তোমাকে ধন্যবাদ".