

Operating System Principles


操作系统原理

Deadlock

李旭东

leexudong@nankai.edu.cn
Nankai University


Case: Traffic



leexudong@nankai.edu.cn

2

Case: Traffic



leexudong@nankai.edu.cn

3

Objectives

- Deadlock Definition
- Deadlock Conditions
- Deadlock Modeling
- Deadlock Detection
- Deadlock Recovery
- Deadlock Avoidance
- Deadlock Prevention

leexudong@nankai.edu.cn

4



Deadlock

- Deadlock
 - A situation in which two or more competing actions are each waiting for the other to finish, and thus neither ever does
 - A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.
 - Deadly-Embrace



The Reason of Deadlock

- Deadlock
 - Competing Resource
 - Running Order of a Set of Processes



Resource

- Resource which causes deadlock
- hardware, software, information
- Preemptable, Nonpreemptable
 - 抢占式、非抢占式



Nonpreemptable Resources

- Sequence of events required to use a resource:
 - Request the resource
 - Use the resource
 - Release the resource

Resource Acquisition

```
typedef int semaphore;
semaphore resource_1;
```

```
void process_A(void) {
    down(&resource_1);
    use_resource_1();
    up(&resource_1);
}
```

(a)

```
typedef int semaphore;
semaphore resource_1;
semaphore resource_2;
```

```
void process_A(void) {
    down(&resource_1);
    down(&resource_2);
    use_both_resources();
    up(&resource_2);
    up(&resource_1);
}
```

(b)

leexudong@nankai.edu.cn

9

Which situation will be a deadlock?

```
semaphore resource_1;
semaphore resource_2;
void process_A(void){
    down(&resource_1);
    down(&resource_2);
    use_both_resources();
    up(&resource_2);
    up(&resource_1);
}
void process_B(void){
    down(&resource_2);
    down(&resource_1);
    use_both_resources();
    up(&resource_1);
    up(&resource_2);
}
```

(a)

```
semaphore resource_1;
semaphore resource_2;
void process_A(void){
    down(&resource_1);
    down(&resource_2);
    use_both_resources();
    up(&resource_2);
    up(&resource_1);
}
void process_B(void){
    down(&resource_1);
    down(&resource_2);
    use_both_resources();
    up(&resource_2);
    up(&resource_1);
}
```


(b)

leexudong@nankai.edu.cn

10

Deadlock Model

SRAG(System Resource Allocation Graph)

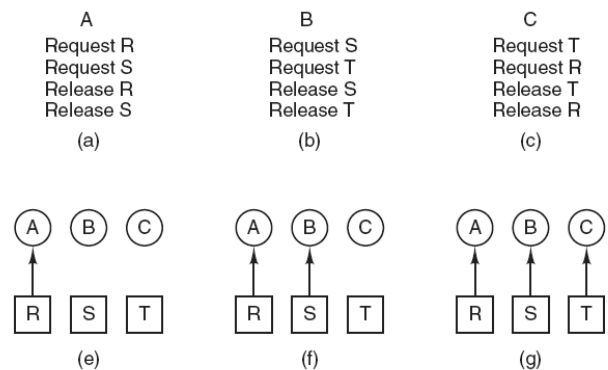


SRAG=(V,E)

leexudong@nankai.edu.cn

11

Deadlock Model






leexudong@nankai.edu.cn




12

Deadlock Model

1. A requests R
2. B requests S
3. C requests T
4. A requests S
5. B requests T
6. C requests R



Deadlock!






leexudong@nankai.edu.cn




13

Deadlock Model

1. A requests R
2. C requests T
3. A requests S
4. C requests R
5. A releases R
6. A releases S



No Deadlock!



leexudong@nankai.edu.cn

14

Necessary Conditions

- Coffman et al., 1971
 - Mutual exclusion
 - Hold and wait
 - No preemption
 - Circular wait

leexudong@nankai.edu.cn

15

deadlock depend on a policy

- ?one resource can be shared among multi-processes
- ?a process which holds a resource requests another resource
- ?policy which allows circular wait
- ?multi-processes concurrent

leexudong@nankai.edu.cn

16

Strategies for Deadlocks


- Strategies for dealing with deadlocks:
 - Detection and recovery
 - let deadlocks occur, detect them, take action.
 - Dynamic avoidance
 - by careful resource allocation
 - Prevention
 - by structurally negating one of the four required conditions
 - ...

leexudong@nankai.edu.cn

17

Strategies for Deadlocks

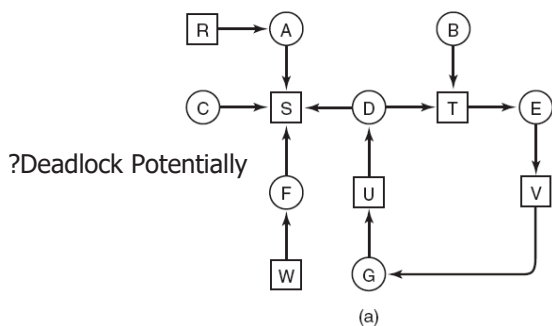
- the ostrich algorithm Just ignore the problem



leexudong@nankai.edu.cn

18

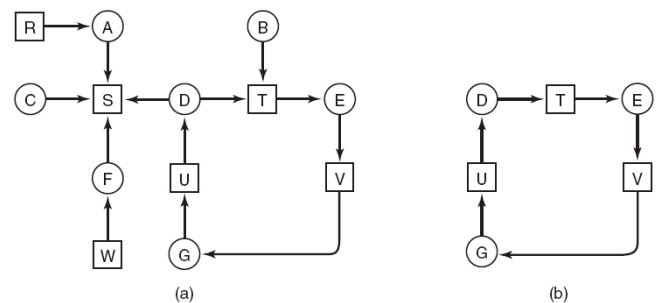
Deadlock Detection with One Resource of Each Type (1)



leexudong@nankai.edu.cn

19

Deadlock Detection with One Resource of Each Type (2)



leexudong@nankai.edu.cn

20

Deadlock Detection with One Resource of Each Type

- Algorithm for detecting deadlock:
 1. For each node, N in the graph, perform the following five steps with N as the starting node.
 2. Initialize L to the empty list, designate all arcs as unmarked.
 3. Add current node to end of L, check to see if node now appears in L two times. If it does, graph contains a cycle (listed in L), **algorithm terminates**.
 - ...

leexudong@nankai.edu.cn

21


Deadlock Detection with One Resource of Each Type

4. From given node, see if any unmarked outgoing arcs. If so, go to step 5; if not, go to step 6.
5. Pick an unmarked outgoing arc at **random and mark** it. Then follow it to the new current node and go to step 3.
6. If this is initial node, graph does not contain any cycles, **algorithm terminates**. Otherwise, dead end. Remove it, go back to previous node, make that one current node, go to step 3.

leexudong@nankai.edu.cn

22

Deadlock Detection with Multiple Resources of Each Type (1)



identical equation:
$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

leexudong@nankai.edu.cn

23

Deadlock Detection with Multiple Resources of Each Type (2)

- Deadlock detection algorithm:
 1. Look for an unmarked process, P_i , for which the i-th row of R is less than or equal to A.
 2. If such a process is found, add the i-th row of C to A, mark the process, and go back to step 1.
 3. If no such process exists, the algorithm terminates.

leexudong@nankai.edu.cn

24

Deadlock Detection with Multiple Resources of Each Type (3)

$$E = \begin{pmatrix} 4 & 2 & 3 & 1 \end{pmatrix} \quad A = \begin{pmatrix} 2 & 1 & 0 & 0 \end{pmatrix}$$

Tape drives Plotters Scanners CD Roms

Current allocation matrix

$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Request matrix

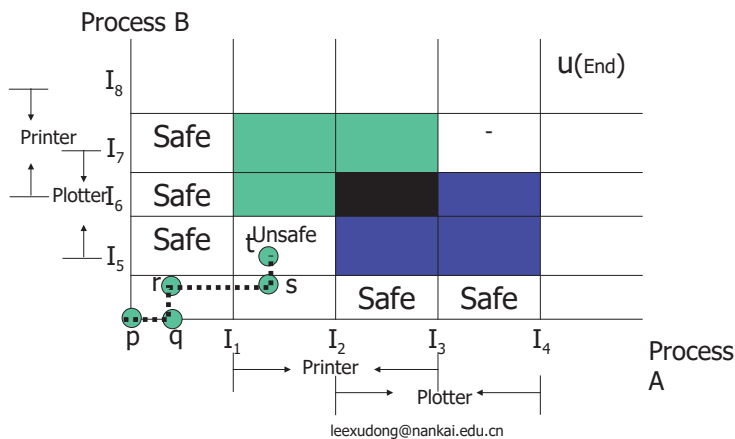
$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

Recovery from Deadlock

- Recovery through preemption
- Recovery through rollback
- Recovery through killing processes

Deadlock Avoidance 避免

Resource Trajectories (资源轨迹图) & Safe State(安全状态)



Safe and Unsafe States (1)

| | Has | Max | |
|-------------|-----|-----|--|
| A | 3 | 9 | |
| B | 2 | 4 | |
| C | 2 | 7 | |
| Free: 3 (a) | | | |

If ...

| | Has | Max | |
|-------------|-----|-----|--|
| A | 4 | 9 | |
| B | 2 | 4 | |
| C | 2 | 7 | |
| Free: 2 (b) | | | |

| | Has | Max | |
|-------------|-----|-----|--|
| A | 4 | 9 | |
| B | 4 | 4 | |
| C | 2 | 7 | |
| Free: 0 (c) | | | |

| | Has | Max | |
|-------------|-----|-----|--|
| A | 4 | 9 | |
| B | — | — | |
| C | 2 | 7 | |
| Free: 4 (d) | | | |

Safe and Unsafe States (2)


| Has Max | | | Has Max | | | Has Max | | |
|---------|---|---|---------|---|---|---------|---|---|
| A | 3 | 9 | A | 3 | 9 | A | 3 | 9 |
| B | 2 | 4 | B | 4 | 4 | B | 0 | — |
| C | 2 | 7 | C | 2 | 7 | C | 2 | 7 |
| Free: 3 | | | Free: 1 | | | Free: 5 | | |
| (a) | | | (b) | | | (c) | | |

| Has Max | | | Has Max | | |
|---------|---|---|---------|---|---|
| A | 3 | 9 | A | 3 | 9 |
| B | 0 | — | B | 0 | — |
| C | 7 | 7 | C | 0 | — |
| Free: 0 | | | Free: 7 | | |
| (d) | | | (e) | | |

leexudong@nankai.edu.cn

29

Safe and Unsafe States (3)



leexudong@nankai.edu.cn

30

The Banker's Algorithm for a Single Resource

| Has Max | | | Has Max | | | Has Max | | |
|----------|---|---|---------|---|---|---------|---|---|
| A | 0 | 6 | A | 1 | 6 | A | 1 | 6 |
| B | 0 | 5 | B | 1 | 5 | B | 2 | 5 |
| C | 0 | 4 | C | 2 | 4 | C | 2 | 4 |
| D | 0 | 7 | D | 4 | 7 | D | 4 | 7 |
| Free: 10 | | | Free: 2 | | | Free: 1 | | |
| (a) | | | (b) | | | (c) | | |

leexudong@nankai.edu.cn

31

The Banker's Algorithm for Multiple Resources (1)

| | Process | Tape drives | Plotters | Printers | CD ROMs |
|--------------------|---------|-------------|----------|----------|---------|
| A | 3 | 0 | 1 | 1 | |
| B | 0 | 1 | 0 | 0 | |
| C | 1 | 1 | 1 | 0 | |
| D | 1 | 1 | 0 | 1 | |
| E | 0 | 0 | 0 | 0 | |
| Resources assigned | | | | | |

| | Process | Tape drives | Plotters | Printers | CD ROMs |
|------------------------|---------|-------------|----------|----------|---------|
| A | 1 | 1 | 0 | 0 | |
| B | 0 | 1 | 1 | 2 | |
| C | 3 | 1 | 0 | 0 | |
| D | 0 | 0 | 1 | 0 | |
| E | 2 | 1 | 1 | 0 | |
| Resources still needed | | | | | |

E = (6342)
P = (5322)
A = (1020)

leexudong@nankai.edu.cn

32

The Banker's Algorithm for Multiple Resources (2)

- Algorithm for checking to see if a state is safe:
 - 1. Look for row, R , whose unmet resource needs all $\leq A$. If no such row exists, system will eventually deadlock since no process can run to completion
 - 2. Assume process of row chosen requests all resources it needs and finishes. Mark process as terminated, add all its resources to the A vector.
 - 3. Repeat steps 1 and 2 until either all processes marked terminated (initial state was safe) or no process left whose resource needs can be met (there is a deadlock).

leexudong@nankai.edu.cn

33

Deadlock Prevention

- Attacking the mutual exclusion condition
- Attacking the hold and wait condition
- Attacking the no preemption condition
- Attacking the circular wait condition


leexudong@nankai.edu.cn

34

Attacking the Circular Wait Condition

1. Imagesetter
2. Scanner
3. Plotter
4. Tape drive
5. CD-ROM drive

(a)



(b)

Approaches to Deadlock Prevention

| Condition | Approach |
|------------------|---------------------------------|
| Mutual exclusion | Spool everything |
| Hold and wait | Request all resources initially |
| No preemption | Take resources away |
| Circular wait | Order resources numerically |

leexudong@nankai.edu.cn

35

leexudong@nankai.edu.cn

36

Order resources numerically

- $R_i (1, \dots, n)$
- $f(R_i)$: the id of R_i
- How to prove
 - If deadlock, so circular wait
 - i.e. existed:
 - $f(R_1) < f(R_2) < \dots < f(R_n) < f(R_1)$
 - **contradiction**

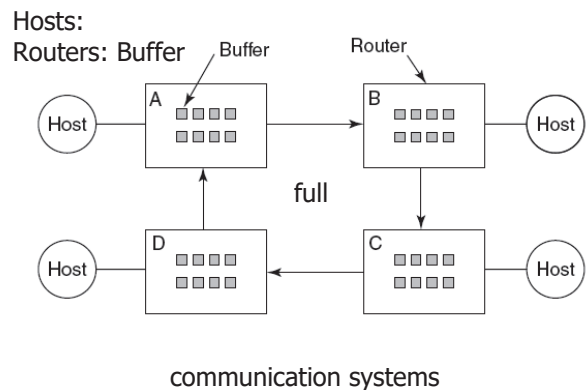
Other Issues

- Two-phase locking
- Communication deadlocks
- Livelock
- Starvation

Two-phase locking

- In the first phase, the process tries to lock all the records it needs, one at a time.
 - If it succeeds, it begins the second phase, performing its updates and releasing the locks. No real work is done in the first phase.
- If during the first phase, some record is needed that is already locked
 - the process just releases all its locks and starts the first phase all over.

Communication deadlocks



Livelock

Case: Each one needs two resources and they use the **polling** primitive **enter_region** to try to acquire the necessary locks. If the attempt fails, the process just **tries again**.

```
void process_A(void) {
    enter_region(&resource_1);
    enter_region(&resource_2);
    use_both_resources( );
    leave_region(&resource_2);
    leave_region(&resource_1);
}

void process_B(void) {
    enter_region(&resource_2);
    enter_region(&resource_1);
    use_both_resources( );
    leave_region(&resource_1);
    leave_region(&resource_2);
}
```

Busy waiting that can lead to livelock

Starvation

Starvation

- In a dynamic system, requests for resources happen all the time.
- Some policy is needed to make a decision about who gets which resource when.
- This policy, although seemingly reasonable, may lead to some processes never getting service even though they are not deadlocked.


Case:

- SJF


Deadlock, livelock, starvation

| | |
|-------------------------|---|
| critical section | A section of code within a process that requires access to shared resources and which may not be executed while another process is in a corresponding section of code. |
| deadlock | A situation in which two or more processes are unable to proceed because each is waiting for one of the others to do something. |
| livelock | A situation in which two or more processes continuously change their state in response to changes in the other process(es) without doing any useful work. |
| mutual exclusion | The requirement that when one process is in a critical section that accesses shared resources, no other process may be in a critical section that accesses any of those shared resources. |
| race condition | A situation in which multiple threads or processes read and write a shared data item and the final result depends on the relative timing of their execution. |
| starvation | A situation in which a runnable process is overlooked indefinitely by the scheduler; although it is able to proceed, it is never chosen. |

?Deadlock



(Optional) Parallel Universes



leexudong@nankai.edu.cn

45

Summary

- Deadlock Definition
- Deadlock Conditions
- Deadlock Modeling
- Deadlock Detection
- Deadlock Recovery
- Deadlock Avoidance
- Deadlock Prevention

leexudong@nankai.edu.cn

46